



JavaScript Spessore

A Thick Shot of Objects, Metaobjects, & Protocols
by Reginald “raganwald” Braithwaite

JavaScript Spessore

A Thick Shot of Objects, Metaobjects, & Protocols

Reginald Braithwaite

This book is for sale at <http://leanpub.com/javascript-spessore>

This version was published on 2013-12-30



This is a [Leanpub](#) book. Leanpub empowers authors and publishers with the Lean Publishing process. [Lean Publishing](#) is the act of publishing an in-progress ebook using lightweight tools and many iterations to get reader feedback, pivot until you have the right book and build traction once you do.

©2013 Reginald Braithwaite

Also By Reginald Braithwaite

Kestrels, Quirky Birds, and Hopeless Egocentricity

What I've Learned From Failure

How to Do What You Love & Earn What You're Worth as a Programmer

CoffeeScript Ristretto

JavaScript Allongé

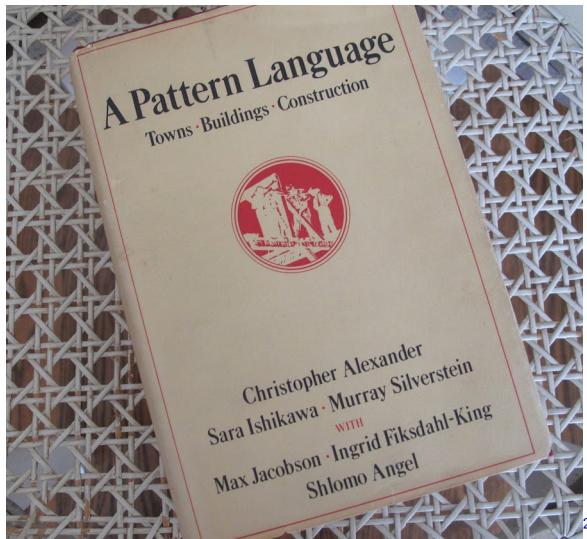
Contents

Prefaces	i
Foreword	ii
Taking a page out of LiSP	v
Disclaimer	vii
What is a Metaobject Protocol?	1
What is an Object?	2
What is a Metaobject?	5
What is a Protocol?	8
Summary	11

Prefaces

Foreword

There are many ways to write books about Architecture. One might attempt to catalogue a broad variety of problems and their solutions, at scales ranging from designing communities down to items of furniture, as Christopher Alexander did in [A Pattern Language](#)¹:



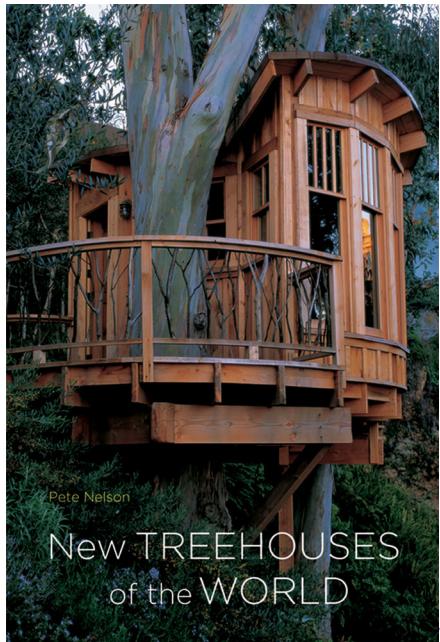
New Treehouses of the World

One might also pick a particular domain and then survey how different architects approached solving a common set of problems, highlighting the variety of possible styles. [New Treehouses of the World](#)³ takes this approach:

¹http://www.amazon.com/gp/product/0195019199/ref=as_li_ss_tl?ie=UTF8&camp=1789&creative=390957&creativeASIN=0195019199&linkCode=as2&tag=raganwald001-20

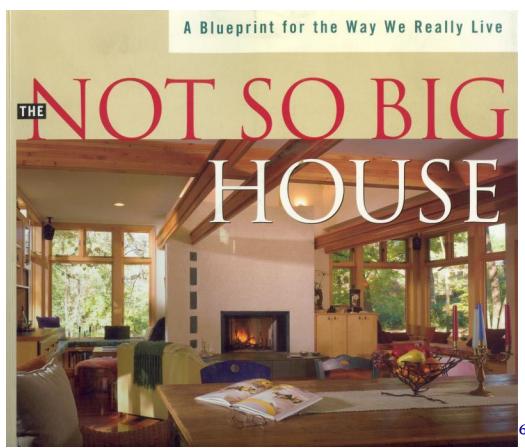
²http://www.amazon.com/gp/product/0195019199/ref=as_li_ss_tl?ie=UTF8&camp=1789&creative=390957&creativeASIN=0195019199&linkCode=as2&tag=raganwald001-20

³http://www.amazon.com/gp/product/0810996324/ref=as_li_ss_tl?ie=UTF8&camp=1789&creative=390957&creativeASIN=0810996324&linkCode=as2&tag=raganwald001-20



The Not So Big House

Or one might approach architecture from a single viewpoint, describing a highly specific and coherent style backed by personal experience, as Sarah Susanka did in [The Not So Big House⁵](#):



⁴http://www.amazon.com/gp/product/0810996324/ref=as_li_ss_tl?ie=UTF8&camp=1789&creative=390957&creativeASIN=0810996324&linkCode=as2&tag=raganwald001-20

⁵http://www.amazon.com/gp/product/1600851509/ref=as_li_ss_tl?ie=UTF8&camp=1789&creative=390957&creativeASIN=1600851509&linkCode=as2&tag=raganwald001-20

⁶http://www.amazon.com/gp/product/1600851509/ref=as_li_ss_tl?ie=UTF8&camp=1789&creative=390957&creativeASIN=1600851509&linkCode=as2&tag=raganwald001-20

JavaScript Spessore

*JavaScript Spessore*⁷ follows this path. This book will not attempt to dictate the “one true definition” of object-oriented programming. It will not survey and catalog the many different programming languages and libraries people have used to solve software problems with objects.

Instead, *JavaScript Spessore* describes one way to design programs with objects as the central idea: Building software on top of a metaobject protocol. A metaobject protocol is a kind of abstraction for implementing object-oriented programming semantics. One program might be developed with pattern matching for methods, another might use state machines, a third might be heavily factored into aspects.

By implementing this metaobject protocol and semantics on top of it, we gain insight into how these semantics work and how they might be applied to solving problems we encounter as software developers.

⁷<https://leanpub.com/javascript-spessore>

Taking a page out of LiSP

Teaching Lisp by implementing Lisp is a long-standing tradition. We read book after book, lecture after lecture, blog post after blog post, all explaining how to implement Lisp in Lisp. Christian Queinnec's [Lisp in Small Pieces](#)⁸ ("LiSP") is particularly notable, not just implementing a Lisp in Lisp, but covering a wide range of different semantics within Lisp.

LiSP's approach is to introduce a feature of Lisp, then develop an implementation. The book covers [Lisp-1 vs. Lisp-2⁹, then discusses how to implement namespaces, building a simple Lisp-1 and a simple Lisp-2. Another chapter discusses scoping, and again you build interpreters for dynamic and block scoped Lisps.

Building interpreters (and eventually compilers) may seem esoteric compared to tutorials demonstrating how to build a blogging engine, but there's a method to this madness. If you implement block scoping in a "toy" language, you gain a deep understanding of how closures really work in any language. You gain some insight into the implications with respect to memory and performance. If you write a Lisp that rewrites function calls in [Continuation Passing Style](#)¹⁰, you can't help but feel comfortable using JavaScript callbacks in [Node.js](#)¹¹.

The simple fact is that *implementing* a language feature teaches you a tremendous amount about how the feature works in a relatively short amount of time. And that goes double for implementing variations on the same feature—like dynamic vs block scoping or single vs multiple namespaces.

That being said, you get the most mileage out of implementing language semantics, not language syntax. Semantics are the *meanings* of programs, syntax is merely the appearance. Writing parsers for both `compose = (a, b) -> (c) -> a(b(c))` in CoffeeScript and `function compose (a, b) { return function (c) { return a(b(c)); }; }` in JavaScript wouldn't teach you nearly as much as writing the code that implements closures. And it's the closures that make `compose` work the way it does.

In this book, we are going to implement a number of different programming language semantics, all in JavaScript. We won't be choosing features at random; We aren't going to try to implement every possible type of programming language semantics. We won't explore dynamic vs block scoping, we won't implement call-by-name, and we will ignore the temptation to experiment with lazy evaluation.

We *are* going to implement different object semantics, implement different kinds of metaobjects, and implement different kinds of method protocols. We are going to focus on the semantics of objects, metaobjects, and protocols, because we're interested in understanding "object-oriented programming" and all of its rich possibilities.

⁸http://www.amazon.com/gp/product/B00AKE1U6O/ref=as_li_ss_tl?ie=UTF8&camp=1789&creative=390957&creativeASIN=B00AKE1U6O&linkCode=as2&tag=raganwald001-20

⁹A "Lisp-1" has a single namespace for both functions and other values. A "Lisp-2" has separate namespaces for functions and other values. To the extend that JavaScript resembles a Lisp, it resembles a Lisp-1. See [The function namespace](#).

¹⁰https://en.wikipedia.org/wiki/Continuation-passing_style

¹¹<http://nodejs.org/about/>

In doing so, we'll learn about the principles of object-oriented programming in far more depth than we would if we chose to implement a "practical" example like a blogging engine.

Disclaimer

Writing is a journey, not a destination. This sample documents the direction we're facing as we take the next step.

This sample “document” is provided to illustrate direction [JavaScript Spessore¹²](#) is taking. **It is not held out to contain any of the actual book’s content.** “Purchases” are being offered to people whose primary motivation is to support the book and encourage me to get it done.

Your purchase does include the right to download any and all versions of the book, as per Leanpub’s lean publishing model. You will never be asked to pay more. You have the right to a 100% no-questions-asked refund if you are not satisfied with the progress of the book:

Within 45 days of purchase you can get a 100% refund on any Leanpub purchase, in two clicks. We process the refunds manually, so they may take a few days to show up. [See full terms¹³](#).

If you buy a Leanpub book you get all the updates to the book for free! All books are available in PDF, EPUB (for iPad) and MOBI (for Kindle). There is no DRM. There is no risk, just guaranteed happiness or your money back.

That being said, writing like this cannot be rushed: progress may be slow relative to a “Learn JavaScript OO in 21 Days” type of book. If you would prefer that there be a substantial amount of work completed, please be patient and check back in a month or so.

Please also bear in mind that pricing and bundling may vary over time. The book may be offered at any price in the future or even be free to read or free to share at some point in the future.

¹²<https://leanpub.com/javascript-spessore>

¹³<https://leanpub.com/terms#returns>

What is a Metaobject Protocol?

A metaobject protocol provides the vocabulary to access and manipulate the structure and behavior of objects. Typical functions of a metaobject protocol include: Creating and deleting new classes; Creating new methods and properties; Changing the class structure so that classes inherit from different classes; Generating or modifying the code that defines the methods for the class.¹⁴[Wikipedia](#)¹⁴

Languages with support for object-oriented programming typically provide their own “Vocabulary to access and manipulate the structure and behavior of objects” at runtime. Ruby, for example, provides keywords like `def` and `class`, as well as imbuing objects and modules with a variety of methods for inspecting and modifying behaviour while the program is running.

JavaScript can be said to have an extremely minimalist metaobject protocol. In this book, we’ll look at building our own, richer metaobject protocols on top of JavaScript. Our metaobject protocols will be guided by the philosophy that object behaviour should be divided amongst three entities:

- Objects are responsible for the domain-specific properties of entities we’re modeling.
- Metaobjects are responsible for encapsulating the behaviour of base objects.
- Protocols are façades for insulating code that defines behaviour from the implementation of metaobjects.

¹⁴<https://en.wikipedia.org/wiki/Metaobject>

What is an Object?

A Smalltalk object can do exactly three things: Hold state (references to other objects), receive a message from itself or another object, and in the course of processing a message, send messages to itself or another object.¹⁵[Smalltalk on Wikipedia](https://en.wikipedia.org/wiki/Smalltalk)¹⁵

Objects seem simple enough: They hold state, they receive messages, they process those messages, and in the course of processing messages, they also send messages. This brief definition implies an important idea: Objects can *change state* in the course of processing messages. They do this directly by removing, changing, or adding to the references they hold.

messages and method invocations



A nine year-old messenger boy

Smalltalk speaks of “messages.” The metaphor of a message is very clear. A message is composed and sent from one entity (the “sender”) to one entity (the “recipient”). The sender specifies the identity of the recipient. The message may contain information for the recipient, instructions to perform some task, or a question to be answered. An immediate reply may be requested, or the sender may trust the message’s recipient to act appropriately. The metaphor of the “message” emphasizes the arms-length relationship between sender and recipient.¹⁶

Popular languages don’t usually discuss messages. Instead, they speak of “invoking methods.” Invoking a method has a much more imperative implication than “sending a message.” It implies that

¹⁵<https://en.wikipedia.org/wiki/Smalltalk>

¹⁶More exotic messaging protocols are possible. Instead of a message being couriered from one entity to another entity, it could be posted on a public or semi-private space where many recipients could view it and decide for themselves whether to respond. Or perhaps there is a dispatching entity that examines each message and decides who ought to respond, much as an operator might direct your call to the right person within an organization.

the entity doing the invoking is causing something to happen, even if the precise implementation is the receiver's responsibility. Most popular languages are synchronous: The code that invokes the method waits for a response.

methods

A method is a kind of recipe for handling a message. In JavaScript, methods are functions. In other languages, methods are a separate kind of thing than functions. In many languages, the namespace for methods is distinct from the namespace for instance variables and other internal references. In JavaScript, methods and internal state are stored together as properties by default.

If we wish to organize methods separately from internal state, we must impose our own structure.

references and state

You can imagine sending a message to a mathematician: "What's the biggest number: one, five, or four?" You can do that in JavaScript:

```
1 Math.max(1, 5, 4)
2 //=> 5
```

Although this is a message, it is unsatisfying to think of `Math` as an object, because it doesn't have any state. Objects were invented fifty years ago¹⁷ to model entities when building simulations. When making a simulation, an essential design technique is to make each entity with its own independent decision-making ability.

Let's say we're modeling traffic. In real life, each car has its own characteristics like maximum speed. Each car has a driver with their own particular style of driving, and of course different cars have different destinations and perhaps different senses of urgency. Each driver independently responds to the local situation around their car. The same goes for traffic lights, roads... Everything has its own independent behavior.

The way to build a simulation is to provide each simulated entity such as the cars, roads, and traffic lights, with their own little programs. You then embed them in a simulated city with a supervisory program doling out events such as rain. Finally, you start the simulation and see what happens.

The key need is to be able to have entities be independent decision-making units that respond to events from outside of themselves. In essence, we're describing computing units. Computation has, at its heart, a program and some kind of storage representing its state. Although impractical, each entity in a simulation could be a Turing Machine with a long tape.

And thus when we think of "objects," we think of independent computing devices, each with their own storage representing their state: **An object is an entity that uses handlers to respond to**

¹⁷<https://en.wikipedia.org/wiki/Simula> "The Simula Programming Language"

messages. It maintains internal state, and its handlers are responsible for querying and/or updating its state.

We'll have a closer look at implementing objects in Javascript a little later in this book.

What is a Metaobject?

In computer science, a metaobject is an object that manipulates, creates, describes, or implements other objects (including itself). The object that the metaobject is about is called the base object. Some information that a metaobject might store is the base object's type, interface, class, methods, attributes, parse tree, etc.¹⁸[Wikipedia¹⁸](https://en.wikipedia.org/wiki/Metaobject)

It is technically possible to write software just using objects. When we need behaviour for an object, we can give it methods by binding functions to keys in the object:

```

1 var sam = {
2   firstName: 'Sam',
3   lastName: 'Lowry',
4   fullName: function () {
5     return this.firstName + " " + this.lastName;
6   },
7   rename: function (first, last) {
8     this.firstName = first;
9     this.lastName = last;
10    return this;
11  }
12 }
```

We call this a “naïve” object. It has state and behaviour, but it lacks division of responsibility between its state and its behaviour.

This lack of separation has two drawbacks. First, it intermingles properties that are part of the model domain (such as `firstName`) with methods (and possibly other properties, although none are shown here) that are part of the implementation domain. Second, when we needed to share common behaviour, we could have objects share common functions, but does it not scale: There's no sense of organization, no clustering of objects and functions that share a common responsibility.

singleton metaclasses

Metaobjects solve this problem by separating the domain-specific properties of objects from their behaviour and implementation-specific properties. In classic JavaScript, we can use a prototype for the behaviour. We call this a “singleton metaobject.”

¹⁸<https://en.wikipedia.org/wiki/Metaobject>

```

1 var singletonMetaobject = {};
2 var sam = Object.create(singletonMetaobject);
3
4 extend(sam, {
5   firstName: 'Sam',
6   lastName: 'Lowry'
7 });
8 extend(singletonMetaobject, {
9   fullName: function () {
10     return this.firstName + " " + this.lastName;
11   },
12   rename: function (first, last) {
13     this.firstName = first;
14     this.lastName = last;
15     return this;
16   }
17 });

```

Notice that a prototype provides value even when we have no intention of sharing behaviour: Its value is that it separates behaviour from our model's domain properties very neatly:

```

1 Object.getOwnPropertyNames(sam)
2 //=> [ 'firstName', 'lastName' ]

```

Insulating our base objects from our metaobjects is important for hiding implementation details, a core “OO” value. When we provide a method like `fullName`, we’re hiding details of how to compute a full name from other entities, all they need to know is the name of the message to send.

hiding implementation details

But we need to hide implementation details within an object as well. Consider logging.¹⁹ If we wish to write to the console every time an object is renamed, we have many choices. Here are three:

First, we can rewrite our function:

¹⁹The canonical example of a cross-cutting concern. We could also consider implementing undo and transactions, they have similar characteristics.

```

1 singletonMetaobject.rename = function (first, last) {
2   console.log(this.fullName() + " is being renamed " + first + " " + last);
3   this.firstName = first;
4   this.lastName = last;
5   return this;
6 };

```

Second, we can use a [method combinator](#)²⁰:

```

1 var logsNameChange = before( function (first, last) {
2   console.log(this.fullName() + " is being renamed " + first + " " + last);
3 });
4
5 singletonMetaobject.rename = logsNameChange( function (first, last) {
6   this.firstName = first;
7   this.lastName = last;
8   return this;
9 });

```

Third, we could use an aspect-oriented programming library like [YouAreDaChef](#)²¹:

```

1 YouAreDaChef(sam).before('rename', function (first, last) {
2   console.log(this.fullName() + " is being renamed " + first + " " + last);
3 });

```

Never mind which choice we should exercise. The important thing is that whatever we do should be hidden from the `sam` object itself! All it “knows” is that it has a prototype. What happens inside that prototype ought to be opaque. Is the `rename` function rewritten? Is it recombined? Is there a `before_rename` property with list of functions to execute? That should be irrelevant to the `sam` object itself.

The basic principle of the metaobject is that we separate the mechanics of behaviour from the domain properties of the base object. Everything else, like how inheritance is implemented, or whether a method is being logged, is hidden from the object.

²⁰<https://github.com/raganwald/method-combinators>

²¹<https://github.com/raganwald/YouAreDaChef>

What is a Protocol?

We've seen that a metaobject separates an object's domain-specific properties from the implementation details of its behaviour. The implementation details might be simple, such as a few methods that are expressed as functions, or complex, such as an inheritance hierarchy with composite methods.

A metaobject *protocol* (or just "protocol") is an interface for managing metaobjects. While the metaobject is responsible for the object's behaviour, the protocol is responsible for creating, changing, and removing behaviour from metaobjects.

protocols and object creation

Previously, we showed how to use a prototype as a singletonMetaobject:

```
1 var singletonMetaobject = {};
2 var sam = Object.create(singletonMetaobject);
3
4 extend(sam, {
5   firstName: 'Sam',
6   lastName: 'Lowry'
7 });
8 extend(singletonMetaobject, {
9   fullName: function () {
10     return this.firstName + " " + this.lastName;
11   },
12   rename: function (first, last) {
13     this.firstName = first;
14     this.lastName = last;
15     return this;
16   }
17});
```

A protocol can hide this implementation detail behind a function:

```

1  function object(superprototype, propertyDescriptors) {
2    if (arguments.length == 0 || typeof(superprototype) == 'undefined') {
3      superprototype = Object.prototype;
4    }
5    var singletonMetaobject = Object.create(superprototype)
6    return Object.create(singletonMetaobject, propertyDescriptors || {});
7  }
8
9  var sam = object();
10 extend(sam, {
11   firstName: 'Sam',
12   lastName: 'Lowry'
13 });

```

Adding methods to the singleton metaobject is also hidden behind a function. We'll elide the details for the moment:

```

1  function extendWithMethods (object, methodDescriptions) {
2    for (var methodName in methodDescriptions) {
3      // ...
4    }
5    return object;
6  }
7
8  extendWithMethods(sam, {
9    fullName: function () {
10      return this.firstName + " " + this.lastName;
11    },
12    rename: function (first, last) {
13      this.firstName = first;
14      this.lastName = last;
15      return this;
16    }
17 });

```

protocols and advice

In the previous section, we gave three different implementations for adding logging to the `rename` method. The base object is insulated from the details by its metaobject, and the code that attaches the logging should likewise be insulated from the implementation choice.

Our protocol provides the insulation for us:

```
1 function method (object, methodName, optionalOverride) {
2   // ...
3 }
4
5 method(sam, 'rename').unshift( function (firstName, lastName) {
6   console.log(this.fullName() + " is being renamed " + firstName + " " + lastName\
7 );
8 });
9
10 sam.rename('Samuel', 'Lowrie')
11 //=> Sam Lowry is being renamed Samuel Lowrie
12 // { firstName: 'Samuel', lastName: 'Lowrie' }
```

The metaobject hides the implementation of “advice” from the base object. However, the implementation is exposed to the code responsible for adding logging to base objects. This is where the metaobject comes in.

To be specific, we meant that a base object should be insulated from the implementation consequences of rewriting a method, applying a combinator to it, or applying “before advice” to it. But we go further: In addition to insulating the base object from the implications of this choice, we must insulate the programmer from needing to think through and possibly re-invent the implementation of any of these three choices.

Summary

Metaobject protocols are vocabularies for accessing and manipulating the structure and behavior of objects. The protocols in this book will be implemented with three main entities:

- *Objects*: Objects are responsible for domain-specific properties.
- *Metaobjects*: Metaobjects are responsible for the insulating objects from the implementation of their behaviour.
- *Protocols*: Protocols are responsible for insulating other code from the implementation of metaobject behaviour.